

VINCENT ZHONG

vincent.zhong@uwaterloo.ca | github.com/vincentzdz | [linkedin.com/in/vincentzdz/](https://www.linkedin.com/in/vincentzdz/) | 613-701-2828

EDUCATION

University of Waterloo <i>Electrical and Computer Engineering, BASc</i>	Sep 2024 – Apr 2029 (Expected) Waterloo, ON
• Computer Science Club Projects Mentor	

EXPERIENCE

Machine Learning Engineer Intern <i>Shopify</i>	Sep 2025 – Dec 2025 Toronto, ON
• Building query rewriting and inference systems for production search infrastructure serving 2M+ monthly queries; optimizing retrieval and query rewriting performance and quality	
• WIP	
Software Engineering Intern <i>Yupp AI</i>	Feb 2025 – May 2025 Mountain View, CA
• Built and deployed embedding service using Nomic Embed for user memory retrieval, handling 200K queries/day on GCP (FastAPI, PyTorch, Docker)	
• Sped up CI/CD pipelines for ML microservices by cutting build time from 22 to 12 min and reducing time-to-merge by 30% (Docker, GitHub Actions, GCP)	
• Deployed vector retrieval system indexing 1M user memory embeddings with pgvector	
Systems Researcher <i>Software Systems at UW Data Systems Group</i>	Nov 2024 – May 2025 Waterloo, ON
• Co-author: "Retrieval with Learned Dense and Sparse Representations Using Anserini" (ACM SIGIR 2025 Resource & Reproducibility Track)	
• Integrated Snowflake's Arctic text embedding models into a retrieval research toolkit as a new dense encoder using ONNX, implementing tokenization and inference to enable embeddings for passage retrieval benchmarks.	
• Added support for a new, efficient sparse encoding model using ONNX to Anserini, boosting query relevance performance by up to 10% on the leading in-domain benchmark.	
• Developed a search and conversational interface using Vercel's AI SDK, allowing users to query dense and sparse indexes using natural language, while also displaying query relevance metrics.	

PROJECTS

Parallel Deep Learning Techniques <i>PyTorch</i>
• Implemented data parallelism from scratch, featuring model replication to GPU instances, data sharding, and gradient synchronization (CPU averaging).
• Designed parallelized forward/backward passes and inter-stage activation/gradient propagation.
• Validated correctness and convergence against PyTorch's <code>nn.DataParallel</code> and sequential execution baselines.

TECHNICAL SKILLS

Languages: Python, Java, C++, TypeScript, SQL, Bash

Frameworks: PyTorch, FastAPI, React, Next.js

Tools: Git, Docker, GitHub Actions, Redis, PostgreSQL, Google Cloud Platform